

# CS 4476 Project 1

Ashwin Rathie

Ashwin.rathie@gatech.edu


arathie6

903281887

# Part 1: Image filtering

<insert visualization of Gaussian kernel from proj1.ipynb here>

```
Success -- kernel values are correct.  
True
```



<Describe your implementation of my\_imfilter() in words.>

Overall, the point of the my\_imfilter() method was to 1) perform the “sliding” of the kernel over the entire image, 2) calculating the new pixel value using convolution, and 3) assigning the pixel value to the new image.


My implementation pads the image with the appropriate kernel dimension//2 on each side of the image.

It then looks at each pixel and each channel of each pixel as well as the surrounding pixels, the “window”, performs the dot product of the kernel and the “window”, sums it up and assigns it to the filtered image.

# Part 1: Image filtering


## Identity filter

<insert the results from proj1\_test\_filtering.ipynb using 1b cat.bmp with the identity filter here>



## Small blur with a box filter


<insert the results from proj1\_test\_filtering.ipynb using 1b cat.bmp with the box filter here>



# Part 1: Image filtering


## Sobel filter

<insert the results from proj1\_test\_filtering.ipynb using 1b\_cat.bmp with the Sobel filter here>



## Discrete Laplacian filter

<insert the results from proj1\_test\_filtering.ipynb using 1b\_cat.bmp with the discrete Laplacian filter here>



# Part 1: Hybrid images

<Describe your implementation of `create_hybrid_image()` here.>

I created the `low_freq` image by calling `my_imfilter()` on `image1` and the kernel. I then created the `high_freq` image by subtracting the result of a call to `my_imfilter()` on `image2` and the kernel from the original image 2 (`orig - low_freq = high_freq`). The hybrid was created by simply adding the `low_freq` and `high_freq` and clipping the values to be between 0 and 1.

**Cat + Dog**

<insert your hybrid image here>



Cutoff frequency: 7

# Part 1: Hybrid images

## Motorcycle + Bicycle

<insert your hybrid image here>



Cutoff frequency: 4

## Plane + Bird

<insert your hybrid image here>



Cutoff frequency: 9

# Part 1: Hybrid images

## Einstein + Marilyn

<insert your hybrid image here>



Cutoff frequency: 4

## Submarine + Fish

<insert your hybrid image here>



Cutoff frequency: 4



# Part 2: Hybrid images with PyTorch

**Cat + Dog**

<insert your hybrid image here>



**Motorcycle + Bicycle**

<insert your hybrid image here>





# Part 2: Hybrid images with PyTorch

**Plane + Bird**

<insert your hybrid image here>



**Einstein + Marilyn**

<insert your hybrid image here>



# Part 2: Hybrid images with PyTorch

## Submarine + Fish

<insert your hybrid image here>



## Part 1 vs. Part 2

<Compare the run-times of Parts 1 and 2 here, as calculated in proj1.ipynb. What can you say about the two methods?>

Part 1 ran in 8.384s while part 2 ran in 1.683s. It seems that the tensor/PyTorch method is significantly faster. This could be due to code optimizations in the PyTorch library (using `conv2d()` instead of our own manual method of convolution), as well as possibly quicker mathematical operations using the Tensor structure rather than `ndarrays`.

# Tests

<Provide a screenshot of the results when you run `pytest tests` on your final code implementation (note: we will re-run these tests).>

```
(proj1) lawn-128-61-67-176:proj1 ashwin$ pytest tests
===== test session starts =====
platform darwin -- Python 3.6.9, pytest-5.0.1, py-1.8.0, pluggy-0.12.0
rootdir: /Users/ashwin/Google Drive/College/Semester 5/compvis_code/proj1
collected 6 items

tests/unit_test.py ..... [100%]

===== 6 passed in 5.38 seconds =====
(proj1) lawn-128-61-67-176:proj1 ashwin$
```

# Conclusions

<Describe what you have learned in this project. Consider questions like how varying the cutoff frequency value or swapping images within a pair influences the resulting hybrid image. Feel free to include any challenges you ran into.>

Through experimentation, it appears that if you want the low frequency image to be more apparent, you should decrease the cutoff frequency (and increase it to have more of the high frequency image). Swapping images within the pair (low\_freq image becomes the high\_freq image and vice versa) also swaps which image appears more prevalent during closer viewing and which is more prevalent during farther viewing. Overall, this was an interesting project and a good intro to PyTorch

One of the significant challenges I ran into was when writing the `get_item()` method. The code comment instructions said to normalize the pixels and transpose the dimensions before transforming the image into a Tensor. Only after much, much debugging, experimentation, and trial and error *did I realize that the `ToTensor()` transformation method transposes and normalizes for you, and that trying to transpose and normalize yourself before calling the transformation actually causes problems.*